

HW4

October 23, 2024

1 STOR 320 Homework 4 Visualization

Please submit the solution to gradescope by 11:59 PM, Oct 24, Thursday.

Name: Ivy Nangalia

PID: 730670491

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from wordcloud import WordCloud
import requests
from bs4 import BeautifulSoup
```

1.1 Problem 1: Visualization of research areas for UNC School of Data Science and Society (25 points)

In this problem, you will use wordcloud to visualize the research areas of faculty members at UNC SDSS

1.1 Use web scraping techniques to collect the research interests of all faculty members in the UNC School of Data Science and Society from the webpage: https://datascience.unc.edu/people/?wpv_view_count=743&wpv-people-category=faculty&wpv_post_search= (10 points)

- Create a `pd.DataFrame` that contains two columns: `faculty_name` and `research_areas`.
- For example, one row of the dataframe is David Adalsteinsson and biology, interface problems, spatial modeling, algorithm development, creating tools for scientific computing and visualization.
- For faculty members without listed research interests, use `np.nan` to represent missing values.

1.2 Notice that there are 4 pages of faculty members in total. Use web scraping to extract the research areas from all 4 pages and store them in the same format as described in 1.1. (2 points)

- Hint: You can use `'https://datascience.unc.edu/people/?wpv_view_count=743&wpv-people-category=' + str(i)` as the URL to access the content on page `i`.

1.3 How many rows do you obtain in the 1.2? How many rows have missing research areas? (2 points)

Hint: You can use `notnull()` or `isnull()` to check if cells within a `DataFrame` have missing data.

1.4 From the DataFrame in 1.2, select rows with research areas. Combine all the research areas into a single string called `text`. Remove the separator characters: `,` `;` and `and`. (6 points)

1.5 Create a wordcloud use function `WordCloud`. Set the title as **Research Areas for UNC SDSS**. Based on your observation, what are the main research topics for UNC SDSS? (5 points)

```
[ ]: #1.1
url_1 = "https://datascience.unc.edu/people/?
        ↪wpv_view_count=743&wpv-people-category=faculty&wpv_post_search="

# Call the GET request
response_1 = requests.get(url_1)

soup = BeautifulSoup(response_1.text, "html.parser")

[ ]: scrape = ".col-md-8 p , h2"

faculty_name_1 = []
research_areas_1 = []

for a in soup.select(scrape):
    a = a.get_text().strip()
    #print(a)
    if "," in a:
        research_areas_1.append(a)
    elif ";" in a:
        research_areas_1.append(a)
    else:
        faculty_name_1.append(a)
        if len(faculty_name_1) != len(research_areas_1) + 1:
            research_areas_1.append(np.NAN)

faculty_1 = pd.DataFrame({"faculty_name": faculty_name_1, "research_areas":
        ↪research_areas_1})
faculty_1
```

```
[ ]:          faculty_name \
0      David Adalsteinsson
1              Stan Ahalt
2              Jay Aikat
3      Amarjit Budhiraja
4      Iain Carmichael
5              Can Chen
6      Anita Crescenzi
7  Snehalkumar 'Neil' Gaikwad
8      Melissa Haendel
9      Thomas Hofweber
```

```

                                research_areas
0  biology, interface problems, spatial modeling,...
1  Signal, image, and video processing, high-perf...
2  Experimental methods and models in networking ...
3  large deviations, stochastic control, stochast...
4  Computational pathology, deep learning for med...
5  control theory, network science, tensor algebr...
6  impact of time constraints and time pressure o...
7  human-AI alignment; AI and decision making; AI...
8                                     NaN
9  metaphysics, the philosophy of language, the f...

```

```
[ ]: #1.2
```

```

faculty_name = [] # initializing the lists outside the loop
research_areas = []

for i in range(1, 5): # not inclusive of upper bound

    url = "https://datascience.unc.edu/people/?
    ↪wpv_view_count=743&wpv-people-category=faculty&wpv_post_search=&wpv_paged="
    ↪+ str(i)
    response = requests.get(url)
    soup = BeautifulSoup(response.text, "html.parser")

    scrape = ".col-md-8 p , h2"

    index: int = 0

    for a in soup.select(scrape):
        a = a.get_text().strip()
        #print(a)
        if "," in a:
            research_areas.append(a)
        elif ";" in a:
            research_areas.append(a)
        else:
            faculty_name.append(a)
            if len(faculty_name) != len(research_areas) + 1:
                research_areas.append(np.NaN)

faculty = pd.DataFrame({"faculty_name": faculty_name, "research_areas":
    ↪research_areas})
faculty

```

```

[ ]:          faculty_name \
0          David Adalsteinsson

```

1 Stan Ahalt
 2 Jay Aikat
 3 Amarjit Budhiraja
 4 Iain Carmichael
 5 Can Chen
 6 Anita Crescenzi
 7 Snehalkumar 'Neil' Gaikwad
 8 Melissa Haendel
 9 Thomas Hofweber
 10 Hsun-Ta Hsu
 11 Dan Kessler
 12 Shahar Kovalsky
 13 Harlin Lee
 14 Youzuo Lin
 15 Yifei Lou
 16 Terry Magnuson
 17 Richard Marks
 18 Steve Marron
 19 Alex McAvoy
 20 Julie McMurtry
 21 Lina Montoya
 22 Santiago Olivella
 23 Courtney Rivard
 24 Rei Sanchez-Arias
 25 Keriayn Smith
 26 Jack Snoeyink
 27 Justin Sola
 28 Matthew G. Springer
 29 Huaxiu Yao
 30 David Yokum
 31 Weitong Zhang
 32 Chudi Zhong

research_areas
 0 biology, interface problems, spatial modeling,...
 1 Signal, image, and video processing, high-perf...
 2 Experimental methods and models in networking ...
 3 large deviations, stochastic control, stochast...
 4 Computational pathology, deep learning for med...
 5 control theory, network science, tensor algebr...
 6 impact of time constraints and time pressure o...
 7 human-AI alignment; AI and decision making; AI...
 8 NaN
 9 metaphysics, the philosophy of language, the f...
 10 NaN
 11 statistical analysis of networks, post-selecti...
 12 optimization, geometry, computer graphics and ...

```

13 Networks, manifolds, optimal transport, noncon...
14 physics-informed machine learning, deep learni...
15 Image processing, sparse signal recovery, nume...
16 mammalian genetics, genomics, chromatin remode...
17 AR/VR, machine learning, context-aware systems
18 statistics, data science and machine learning,...
19 evolutionary dynamics, game theory, multi-agen...
20 NaN
21 causal inference, precision health/policy, (op...
22 statistical models, particularly Bayesian grap...
23 rhetoric, composition, digital humanities
24 Data mining, machine learning algorithm develo...
25 RNA biology, noncoding RNAs, gene editing, bio...
26 Computational geometry; algorithms for geograp...
27 gun ownership, trends in social research, the ...
28 NaN
29 NaN
30 NaN
31 reinforcement learning, machine learning
32 machine learning, optimization, human-model in...

```

```

[ ]: #1.3
missing = faculty["research_areas"].isnull().sum()

len(faculty), missing

```

```
[ ]: (33, 6)
```

There are 33 rows in the DataFrame from 1.2 in total, and 6 of those rows have null values.

```

[ ]: #1.4
research_areas = faculty["research_areas"].dropna()
text = ", ".join(research_areas).replace(",", " ").replace(";", " ").replace("↵",
↵and", " ")
text[0:200] # sliced to truncate output

```

```
[ ]: 'biology interface problems spatial modeling algorithm development creating
tools for scientific computing visualization Signal image video processing high-
performance scientific industrial computing p'
```

```

[ ]: wordcloud = WordCloud(background_color="white")
wordcloud.generate(text)
plt.imshow(wordcloud)
plt.title("Research Areas for UNC SDSS")
plt.axis("off")
plt.show()

```

[illegible]

1.2 Problem 2. (20 points).


- 2.1 visualize this function for $x \in [-2, 2]$ using `plt`. Add a title to the plot. (5 points)
- 2.2 Add a horizontal line at $y = 0$ to the graph in 2.1 (2 points)

- 2.3 How many intersections of the function and the horizontal line do you have? (3 points)
- 2.3 Based on the plot in 2.2, create another plot that zooms in on the area where x is between $[0.6, 0.7]$. Add a title to the plot. (5 points)
- 2.4 Create another plot to find the intersections of the function and the horizontal line. Mark them with red dots. You may find the intersections by observation. (5 points)

6

```
# 2.2
plt.hlines(0, -2, 2, colors="red", linestyle="dashed")

plt.show()
```

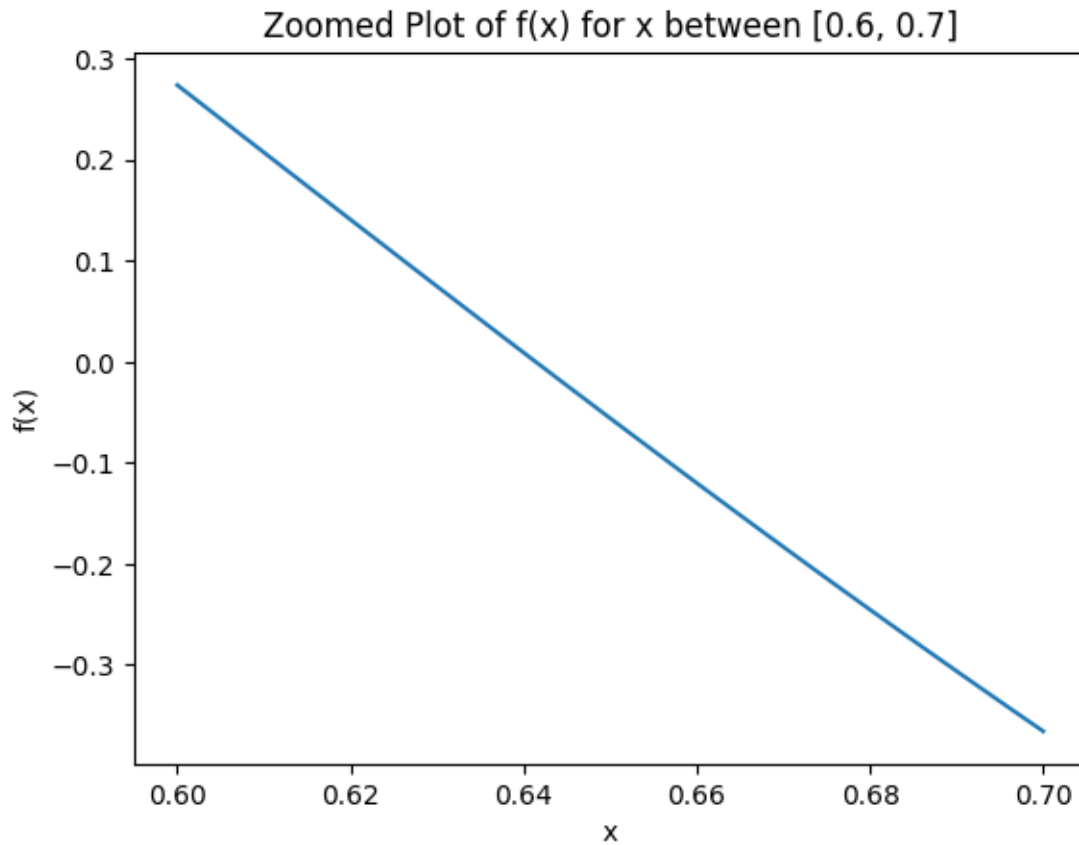


1.2.1 2.3

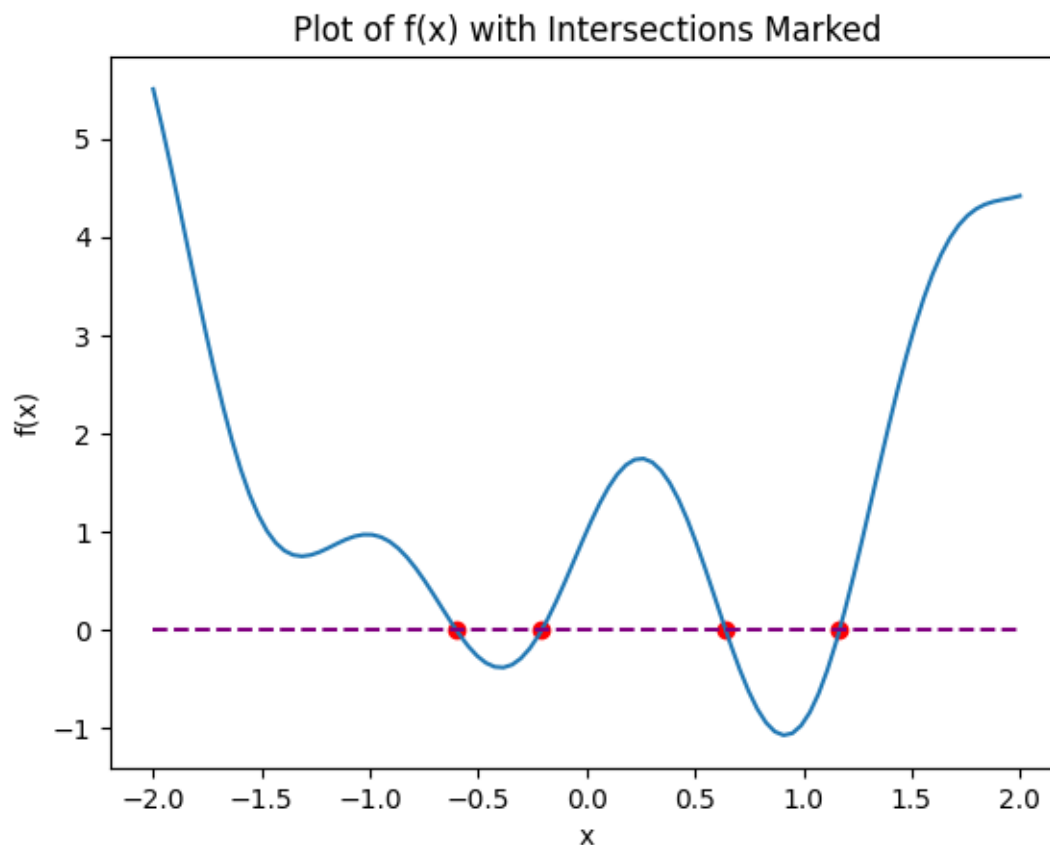
There are four intersections of $f(x)$ and the horizontal line at $x = 0$.

```
[ ]: #2.3
x_zoom = np.linspace(0.6, 0.7, 100)
y_zoom = f(x_zoom)

plt.plot(x_zoom, y_zoom)
plt.title("Zoomed Plot of f(x) for x between [0.6, 0.7]")
plt.xlabel("x")
plt.ylabel("f(x)")
plt.show()
```



```
[ ]: # 2.4
plt.plot(x, y)
plt.hlines(0, -2, 2, colors="purple", linestyle="dashed")
plt.title("Plot of  $f(x)$  with Intersections Marked")
plt.xlabel("x")
plt.ylabel("f(x)")
plt.scatter(1.16, 0, color="red")
plt.scatter(0.64, 0, color="red")
plt.scatter(-0.21, 0, color="red")
plt.scatter(-0.6, 0, color="red")
plt.show()
```

1.3 Problem 3. (20 points).

In this problem, you will draw a 3D ellipsoid in Matplotlib. To do so, you need to parameterize the ellipsoid equation in terms of two angles (azimuthal angle θ and polar angle ϕ and use the parametric equation of the ellipsoid.

The general equation for an ellipsoid is:

$$\frac{x^2}{a^2} + \frac{y^2}{b^2} + \frac{z^2}{c^2} = 1,$$

where: a, b, c are the semi-axes of the ellipsoid along the x , y , and z directions, respectively.

To plot the ellipsoid, you can parametrize it in spherical coordinates as follows:

$$\begin{aligned} x &= a \sin(\phi) \cos(\theta) \\ y &= b \sin(\phi) \sin(\theta) \\ z &= c \cos(\phi), \end{aligned}$$

where: θ varies from 0 to 2π , ϕ varies from 0 to π .

Let $a = 2, b = 1, c = 1.5$. Make a 3D plot that - show the ellipsoid (10 points) - add appropriate title and labels for each axis (5 points) - add two planar cross-sections of the ellipsoid: one parallel to the XY-plane (at $z = 0$) and another parallel to the XZ-plane (at $y = 0$). Use dashed lines to indicate these cross-sections. (5 points)

```
[ ]: ax = plt.axes(projection="3d")

a = 2
b = 1
c = 1.5

theta = np.linspace(0, 2*np.pi, 100)
phi = np.linspace(0, np.pi, 100)

theta, phi = np.meshgrid(theta, phi)


x = a * np.sin(phi) * np.cos(theta)
y = b * np.sin(phi) * np.sin(theta)
z = c * np.cos(phi)

fig = plt.figure()
ax = fig.add_subplot(111, projection="3d")
ax.plot_surface(x, y, z, alpha=0.5)


x_cross = np.linspace(-a, a, 100)
y_cross = np.linspace(-b, b, 100)
z_cross = np.zeros_like(x_cross)

ax.plot(x_cross, y_cross, z_cross, linestyle="dashed", color="red")
ax.plot(x_cross, z_cross, y_cross, linestyle="dashed", color="blue")
ax.set_zlabel("Z")
ax.set_title("3D Ellipsoid")

plt.draw()
plt.show()
```



3D Ellipsoid



1.4 Problem 4 (20 points): You are provided with a dataset containing random information about 50 different cars. The dataset includes the following columns:

- Horsepower: The horsepower of the car.
- Weight: The weight of the car in pounds.
- MPG: The miles per gallon (fuel efficiency) of the car.
- ModelYear: The model year of the car.
- Brand

```
[ ]: cars = {
    'Horsepower': np.random.randint(60, 200, 50),
    'Weight': np.random.randint(1500, 4000, 50),
    'MPG': np.random.uniform(10, 40, 50),
    'ModelYear': np.random.randint(1970, 2020, 50),
    'Brand': np.random.choice(['Toyota', 'Honda', 'KIA'], 50)
}
cars = pd.DataFrame(cars)
cars.head(3)
```

```
[ ]:      Horsepower  Weight      MPG  ModelYear  Brand
0          155     3752  27.942459      1996  Toyota
1          156     3725  29.392193      1982   Honda
2           60     3540  33.480513      2000  Toyota
```

Create a scatter plot to include the information of all five columns in `cars`.

- Use the shape of markers to represent car brands (3 points)
- Use Color to represent the weight of the car (3 points)
- Use x-axis to represent Horsepower (3 points)
- Use y-axis to represent MPG (3 points)
- Use the size of markers to represent the ModelYear (3 points)

Add titles, legend, labels to the graph. (5 points)

```
[ ]: marker_shapes = {"Toyota": "*", "Honda": "s", "KIA": "D"}

for brand, marker in marker_shapes.items():
    brand_data = cars[cars["Brand"] == brand]
    plt.scatter(
        x=brand_data["Horsepower"],
        y=brand_data["MPG"],
        c=brand_data["Weight"],
        s= 3 * (brand_data["ModelYear"] - (cars["ModelYear"].min()-1)),
        marker=marker,
```

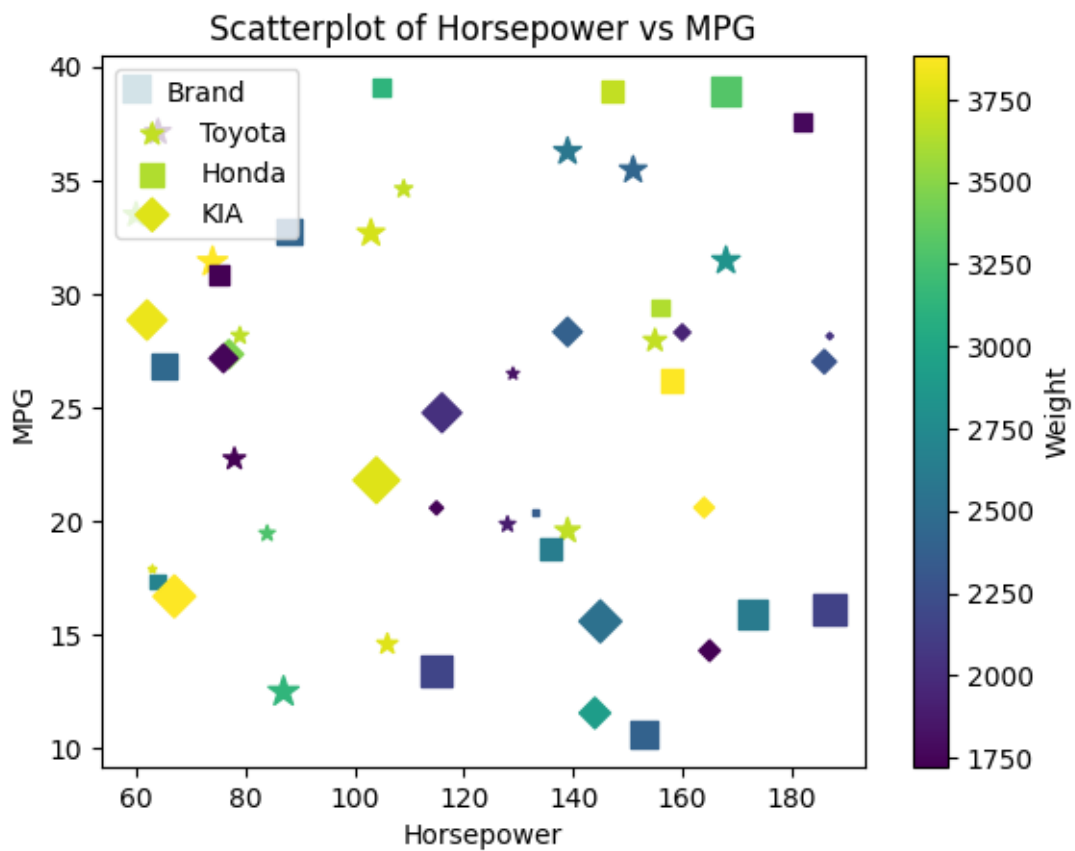
```

        #cmap="viridis",
        #alpha=0.6,
        label=brand
    )

plt.title("Scatterplot of Horsepower vs MPG")
plt.xlabel("Horsepower")
plt.ylabel("MPG")
plt.colorbar(label="Weight")
plt.legend(marker_shapes.keys(), title="Brand", loc="upper left")

plt.show()

```



1.5 Problem 5: Recreate the same plot using sns (15 points)

Recall that in the lecture, we use the following code to create a bar plot for the iris dataset

```

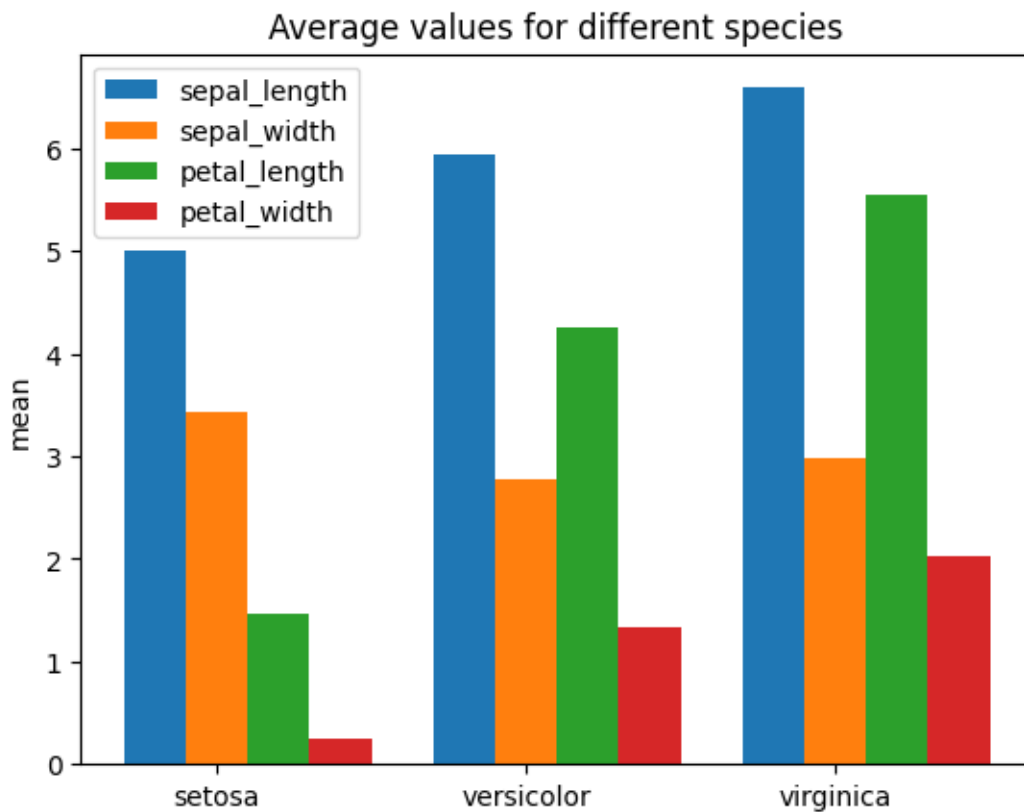
[ ]: df = pd.read_csv('iris.csv')
      avg_iris = df.groupby("species").mean()
      width=0.2

```

```

multiplier=0
x = np.arange(3)
for i in range(avg_iris.shape[1]):
    offset = width*multiplier
    plt.bar(x + offset, avg_iris.iloc[:,i], width, label=avg_iris.columns[i])
    multiplier += 1
plt.legend()
plt.ylabel('mean')
plt.title('Average values for different species')
plt.xticks(x + width, avg_iris.index);

```



Create the same plot using `sns.barplot` function

```

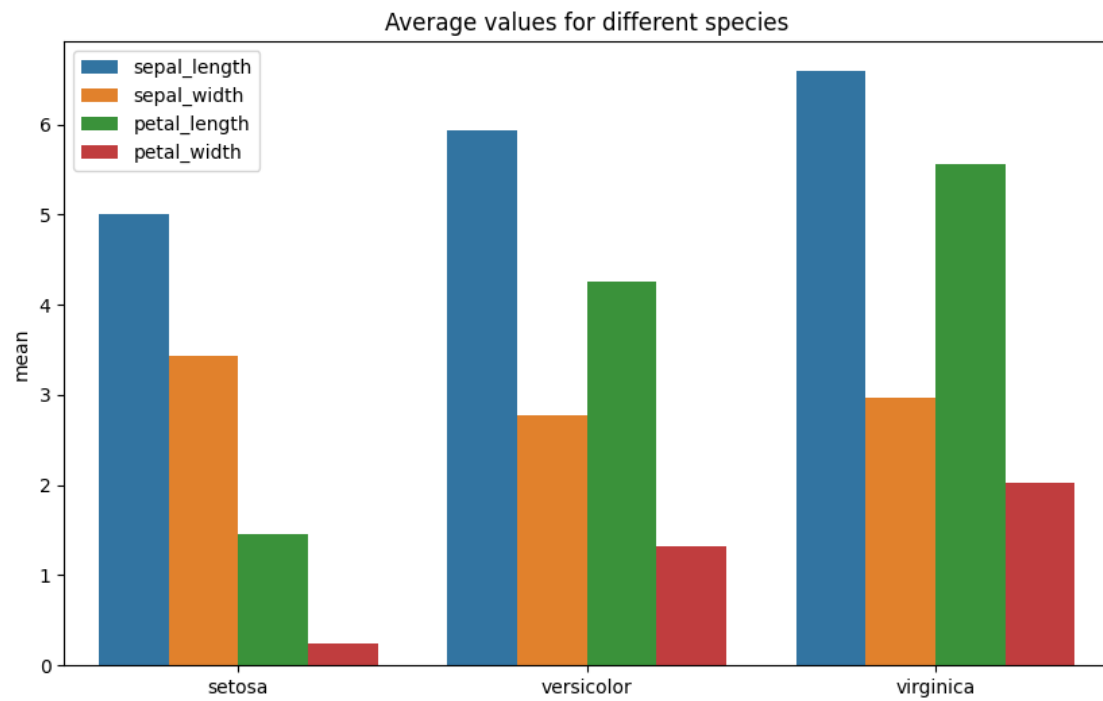
[ ]: avg_iris_long = avg_iris.reset_index().melt(id_vars="species",
    var_name="feature", value_name="mean")

plt.figure(figsize=(10,6))
sns.barplot(data=avg_iris_long, x="species", y="mean", hue="feature")

plt.title("Average values for different species")

```

```
plt.ylabel("mean")  
plt.xlabel("")  
plt.legend()  
plt.show()
```



```
[ ]:
```