



Cross Validation

STOR 320-002 Introduction to Data Science
Fall 2024 - Mo Liu



Review from last lecture

Questions:

- In PCA, when the number of PC gets larger, does the model become more flexible or less flexible?
- In LASSO or Ridge regression, when λ gets larger, does the model become more flexible or less flexible?

+ Hyperparameters of models


3

- PCA: Number of PC
- LASSO or Ridge: scale of regularization parameter


+ Overfitting

4


- Overfitting occurs when the estimated model fits the noise in the training data
- All statistical learning methods are at risk for overfitting



High bias
(underfit)



"Just right"



High variance
(overfit)

+ Example of LASSO

5

- What is the R-squared (R^2) and out-of-sample R-squared (OSR^2) when λ equals the following values?

- 10^{-6}

- 10^{-5}

- 10^{-4}

- 10^{-3}

- 10^{-2}

- 10^{-1}

+ Overfitting

- Tell-tale sign of overfitting is when a model performs very well on the training set in a way that will not translate to good performance on new data
- Overfitted models are often more complicated and do not necessarily accord with intuition/judgment
- Overfitted models often capture idiosyncrasies or errors in the data instead of capturing actual relationships that will hold for new data

+ Overfitting

- Overfitting is possible due to differences between how models are fitted and how they are used
 - Model fitting seeks to maximize accuracy of the model on a known set of training data
 - Models are used to make predictions on new, previously unseen data
- Care must be taken to make sure that the model we estimate does not suffer from overfitting
- How to find the best value of hyperparameters?

+ Grid Search for λ

8

- Generate a candidate set of λ value for LASSO: exponential of $\{-5, -4, -3, -2, -1, 0\}$
- Evaluate the performance at each candidate λ value
 - How to evaluate the performance for λ ?
 - Can we directly look at OSR2?



Training, Test, and Validation Sets

- Recall our usual split into training and test set:
 - Training set is used to build the model
 - Test set is used to assess its performance
- Question: can we estimate test set aka out of sample performance during the training phase *without* touching the test set data?
- Answer: yes! The simplest way to is to create a “virtual” test set called the validation set


+ Validation Set Approach

10

- Step 1: Split into 3 sets – training, validation, and test
- Step 2: Build a sequence of models using the training data (for example, using different λ values in LASSO)
- Step 3: Evaluate each model's performance (e.g. R^2 , accuracy, TPR, FPR, ...) on the validation set
- Step 4: Pick the best model and use the test set to estimate future real-world performance of the model
- (By the way, are you “allowed” to iterate steps 2 and 3?)

+ Training, Test, and Validation Sets

11



Model Building Data

**Model Assessment
Data**

+ Validation Set Approach

- You've probably already sort of used the validation set approach
- For example, if you've looked at the two or more OSR2 values (however that was test data)
- Best practice is to keep the test set in a vault, although it's okay to look at the performance of a handful of models
- This becomes dangerous when we want to check many different models

+ Advantages and Disadvantages of the Validation Set Approach

■ Advantages:

- Conceptually simple
- Easy to implement

■ Disadvantages

- Estimates tend to be sensitive to the results of the random number generator
- “Waste of data” – especially troublesome if n is small

+ Leave-One-Out Cross-Validation (LOOCV)


- As an example, suppose that we want to estimate the error rate (1 minus the accuracy) of a CART classification model
- Idea:
 - Take the training set of size n and remove observation i
 - Build the model on the remaining $n - 1$ observations
 - Check if we made an accurate prediction \hat{y}_i on the i th example that we held out:

$$\text{Err}_i = \begin{cases} 1 & \hat{y}_i \neq y_i \\ 0 & \hat{y}_i = y_i \end{cases}$$

+ Leave-One-Out Cross-Validation (LOOCV)

- LOOCV says to repeat the removal, training process and error calculation for **all** n data samples and use their average to estimate the out of sample performance:

$$CV_{(n)} = \frac{1}{n} \sum_{i=1}^n \text{Err}_i$$



+ Leave-One-Out Cross-Validation (LOOCV)

$$CV_{(n)} = \frac{1}{n} \sum_{i=1}^n \text{Err}_i$$

- Repeat the process of calculating $CV_{(n)}$ for each potential model
- Compare the results as before and pick the best model

+ Advantages and Disadvantages of LOOCV

■ Advantages:

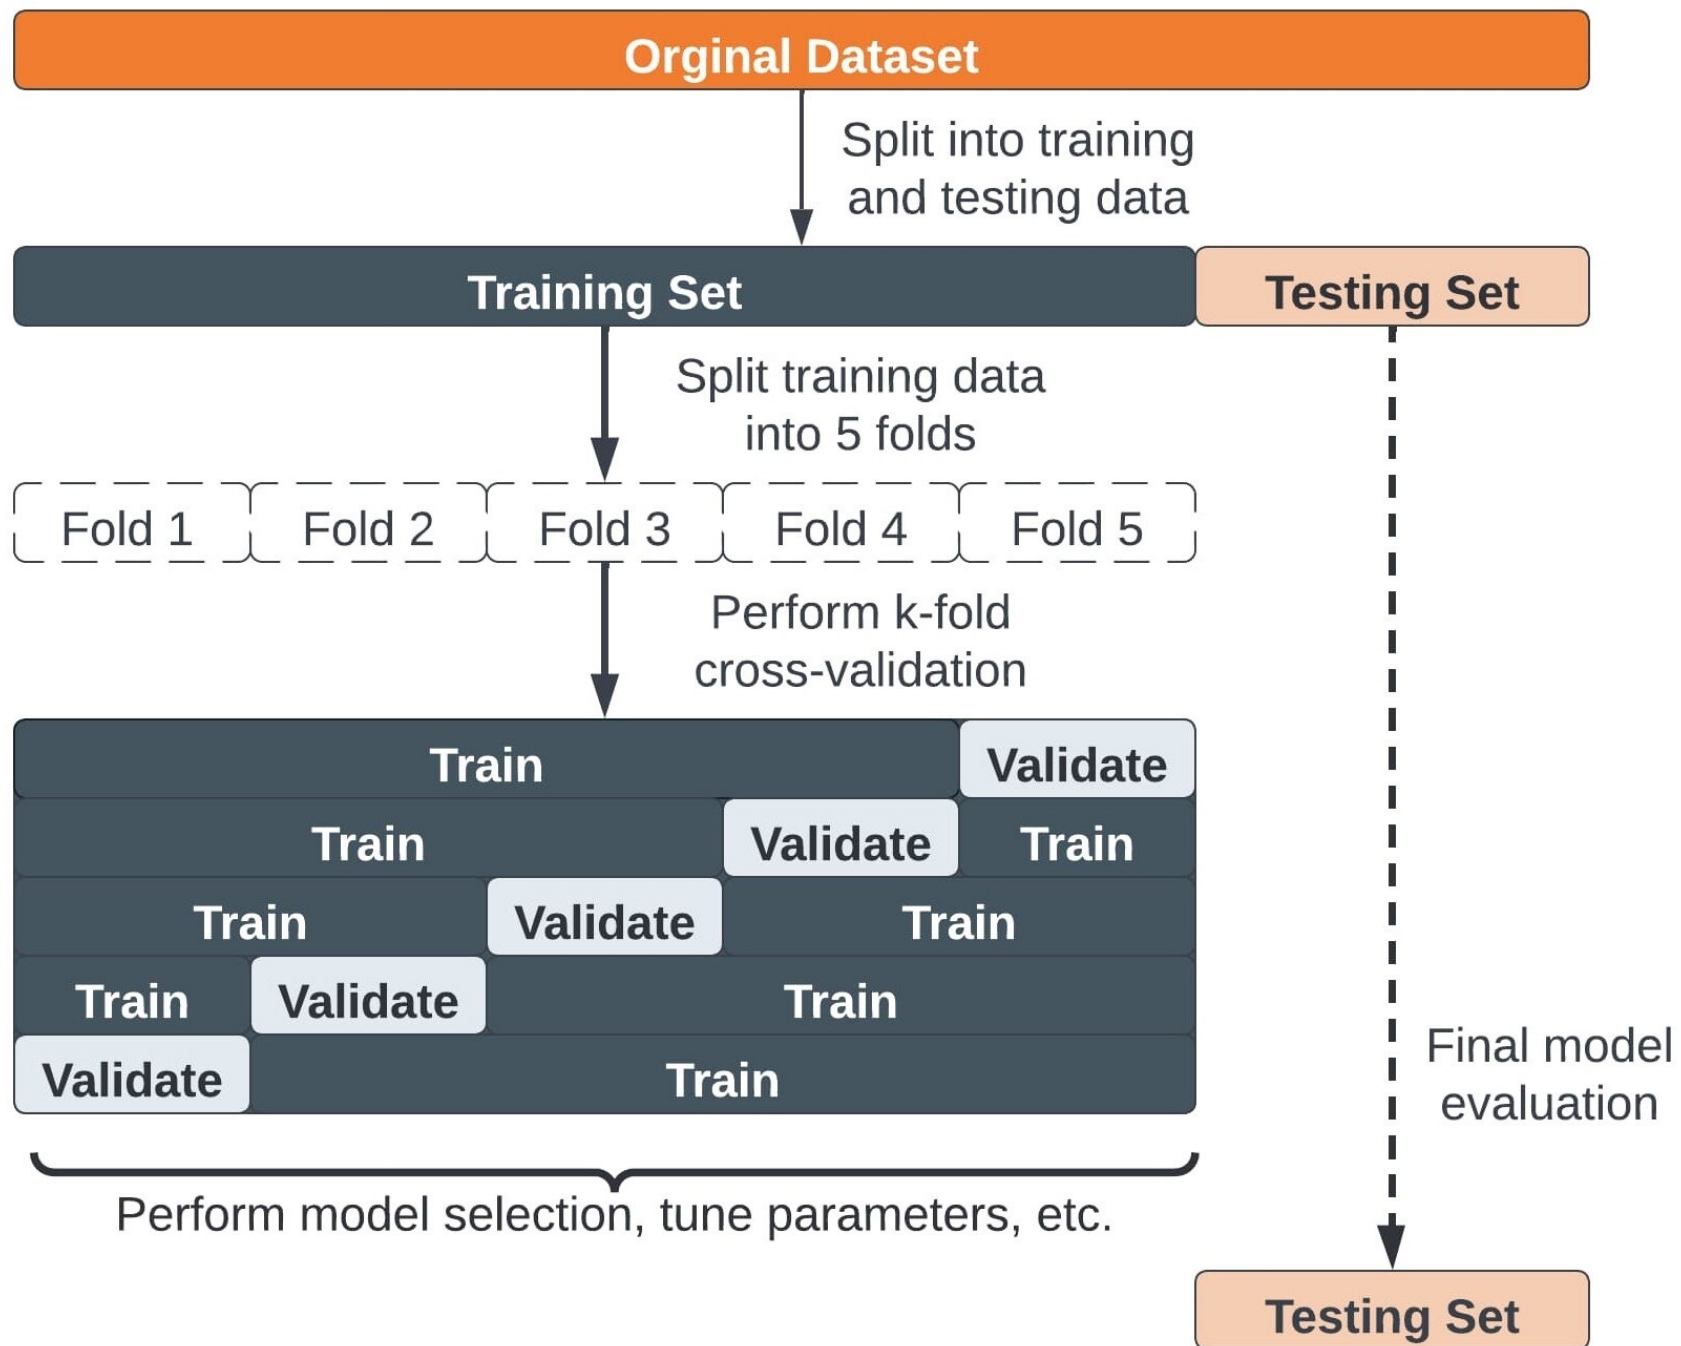
- The result is entirely deterministic given the training data (no random number generation)
- Estimates tends to be more accurate “on average” (less bias)

■ Disadvantages

- Computationally intensive: requires training $n * (\text{number of candidate models})$ different models
- Results within each $CV_{(n)}$ calculation are correlated – averaging correlated numbers may not reduce variance

+ A Compromise: k-fold Cross-Validation

- Step 1: Divide the training data into k different groups (e.g., k = 5 or 10)
- Step 2:
 - Remove group 1 and train on the remaining k-1 groups
 - Compute the average error on the held out group 1, Err_1
 - Repeat for group 2, group 3, ..., group k
- Step 3:
 - Average the results:
$$\text{CV}_{(k)} = \frac{1}{k} \sum_{j=1}^k \text{Err}_j$$
- Repeat steps 2-3 for each candidate model



+ k-fold Cross-Validation

20

- Final step: After selecting your parameters via k-fold cross validation, retrain on the entire dataset using those parameters
- Some more details:
 - $k = 5$ or 10 are commonly used and tend to be good choices
 - For regression, use the MSE (aka RSS) or R^2 as the performance metric
- When $k = 1$, it becomes LOOCV

+ Which methods should we use?

- Validation set approach is better when:
 - The sample size is small
 - The model is time-consuming to build
 - For example, deep learning model
- K-fold cross validation is better when:
 - The sample size is large
 - The model is easy to build

+ k-fold Cross-Validation in Python

- `from sklearn.model_selection import GridSearchCV`
- `alpha_grid = {'alpha': np.logspace(-8, -1, num=50, base=10)}`
- `lasso_cv = GridSearchCV(model, param_grid =
alpha_grid, scoring='neg_mean_squared_error', cv=10,
verbose=1)`
- `lasso_cv.fit(X_train_lasso, y_train)`

+ GridSearchCV, scoring

- GridSearchCV(model, param_grid =, scoring=, cv=10, verbose=1)

Regression		
'explained_variance'	<code>metrics.explained_variance_score</code>	
'neg_max_error'	<code>metrics.max_error</code>	
'neg_mean_absolute_error'	<code>metrics.mean_absolute_error</code>	
'neg_mean_squared_error'	<code>metrics.mean_squared_error</code>	
'neg_root_mean_squared_error'	<code>metrics.root_mean_squared_error</code>	
'neg_mean_squared_log_error'	<code>metrics.mean_squared_log_error</code>	
'neg_root_mean_squared_log_error'	<code>metrics.root_mean_squared_log_error</code>	
'neg_median_absolute_error'	<code>metrics.median_absolute_error</code>	
'r2'	<code>metrics.r2_score</code>	
'neg_mean_poisson_deviance'	<code>metrics.mean_poisson_deviance</code>	
'neg_mean_gamma_deviance'	<code>metrics.mean_gamma_deviance</code>	
'neg_mean_absolute_percentage_error'	<code>metrics.mean_absolute_percentage_error</code>	
'd2_absolute_error_score'	<code>metrics.d2_absolute_error_score</code>	

+ GridSearchCV, cv

- `GridSearchCV(model, param_grid =, scoring='neg_mean_squared_error', cv=10, verbose=1)`

cv : *int, cross-validation generator or an iterable, default=None*

Determines the cross-validation splitting strategy. Possible inputs for cv are:

- None, to use the default 5-fold cross validation,
- integer, to specify the number of folds in a `(Stratified)KFold`,
- [CV splitter](#),
- An iterable yielding (train, test) splits as arrays of indices.

For integer/None inputs, if the estimator is a classifier and `y` is either binary or multiclass, `StratifiedKFold` is used. In all other cases, `KFold` is used. These splitters are instantiated with `shuffle=False` so the splits will be the same across calls.

Refer [User Guide](#) for the various cross-validation strategies that can be used here.

+ Randomly split k-fold sets

- By default, there is no shuffling in Kfold

`sklearn.model_selection.KFold`

```
class sklearn.model_selection.KFold(n_splits=5, *, shuffle=False, random_state=None)
```

[\[source\]](#)

K-Folds cross-validator

Provides train/test indices to split data in train/test sets. Split dataset into k consecutive folds (without shuffling by default).

Each fold is then used once as a validation while the k - 1 remaining folds form the training set.

Read more in the [User Guide](#).

- What if the training set is sorted in some way, for example, by time?




- `from sklearn.model_selection import GridSearchCV`
- `from sklearn.model_selection import KFold`
- `alpha_grid = {'alpha': np.logspace(-8, -1, num=15, base=10)}`
- `cv = KFold(n_splits = 10, random_state = 1, shuffle = True)`
- `lasso_cv = GridSearchCV(lasso, param_grid = alpha_grid, scoring='neg_mean_squared_error', cv=cv, verbose=1)`
- `lasso_cv.fit(X_train_lasso, y_train)`

+ Selecting Parameters via Cross-Validation

■ One Standard Error Rule

- Find model with minimum error
- Select the simplest model whose mean falls within 1 standard deviation of the minimum



+ Custom loss function

- MAE: mean absolute error
- What if we only concern about the case when the prediction error is larger than \$2,000?
- In other words, we want to minimize the number of predictions whose error is larger than \$2,000?

+ Acknowledgement

- The figure of 5-fold cross validation is taken from <https://www.aptech.com/blog/understanding-cross-validation/>