

logistic_regression

November 17, 2024

1 STOR 320 Intro to Data Science: Logistic Regression

- 1. Logistic Regression with numerical variables
- 2. Logistic Regression with categorical variables
- 3. LDA

```
[ ]: import numpy as np
import pandas as pd
```

As usual, we summon `numpy` and `pandas` for dataset representation and manipulation.

1.1 1. LOGISTIC REGRESSION (ONLY NUMERICAL VARIABLES)

1.1.1 1.1 Data loading

```
[ ]: loans = pd.read_csv("loans.csv")
loans.info()
loans.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9516 entries, 0 to 9515
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   not.fully.paid         9516 non-null   int64
1   installment            9516 non-null   float64
2   log.annual.inc         9516 non-null   float64
3   fico                   9516 non-null   int64
4   revol.bal              9516 non-null   float64
5   inq.last.6mths         9516 non-null   int64
6   pub.rec                9516 non-null   int64
dtypes: float64(3), int64(4)
memory usage: 520.5 KB
```

```
[ ]: not.fully.paid  installment  log.annual.inc  fico  revol.bal  \
0                0      829.10      4.929419   737    28.854
1                0      228.22      4.812913   707    33.623
2                0      366.86      4.505150   682     3.511
3                0      162.34      4.929419   712    33.667
```

4	0	102.92	4.907411	667	4.740
---	---	--------	----------	-----	-------

	inq.last.6mths	pub.rec
0	0	0
1	0	0
2	1	0
3	1	0
4	0	0

1.1.2 1.2 Data cleaning and transformation

In Python, a convention is to name variables with underscores. This is slightly different from R. Let us practice how to rename columns in Pandas.

```
[ ]: new_column_names = {'not.fully.paid': 'not_fully_paid', 'log.annual.inc':
    ↪ 'log_annual_inc',
    'revol.bal': 'revol_bal', 'inq.last.6mths': 'inq_last_6mths',
    ↪ 'pub.rec': 'pub_rec'}
loans.rename(columns = new_column_names, inplace = True)

print(loans.columns)
```

```
Index(['not_fully_paid', 'installment', 'log_annual_inc', 'fico', 'revol_bal',
      'inq_last_6mths', 'pub_rec'],
      dtype='object')
```

Use `df.describe()`, you can have a quick overview of the data set. Here, observe that the first four variables ('installment', 'log_annual_inc', 'fico', and 'revol_bal') takes continuous numeric values and the last two variables ('inq_last_6mths', 'pub_rec') takes integer numeric values.

```
[ ]: loans.describe()
```

```
[ ]:      not_fully_paid  installment  log_annual_inc    fico  revol_bal \
count      9516.000000    9516.000000    9516.000000    9516.000000    9516.000000
mean         0.159836      320.131185         4.748642     710.841950      16.988484
std          0.366473      207.069870         0.265002      37.956246     33.721379
min          0.000000       15.670000         3.277838     612.000000       0.000000
25%          0.000000      164.020000         4.588821     682.000000       3.272750
50%          0.000000      269.545000         4.748188     707.000000       8.687500
75%          0.000000      435.405000         4.903323     737.000000      18.354250
max          1.000000      940.140000         6.309584     827.000000     1207.359000

      inq_last_6mths    pub_rec
count      9516.000000    9516.000000
mean         1.572930      0.062211
std          2.200329      0.262406
min          0.000000      0.000000
25%          0.000000      0.000000
```

50%	1.000000	0.000000
75%	2.000000	0.000000
max	33.000000	5.000000

1.1.3 1.3 Train-test split

```
[ ]: from sklearn.model_selection import train_test_split

loans_train, loans_test = train_test_split(loans, test_size=0.3,
↳ random_state=88)
loans_train.shape, loans_test.shape
```

```
[ ]: ((6661, 7), (2855, 7))
```

```
[ ]: loans_train.head()
```

```
[ ]:      not_fully_paid  installment  log_annual_inc  fico  revol_bal  \
6330                0         682.74         5.079181   697      72.682
3799                0         479.21         4.778151   677      14.180
4019                0         333.15         4.579784   702       5.520
7475                0         373.32         4.819544   782       6.449
8641                0          81.51         4.417638   662       0.558

      inq_last_6mths  pub_rec
6330                1         0
3799                1         0
4019                1         0
7475                2         0
8641                4         0
```

test_size = 0.3 means that we will put 30% of the data in the test set, 70% in the training set.

1.1.4 1.4 Baseline model

```
[ ]: # How many loans have defaulted in the training set?

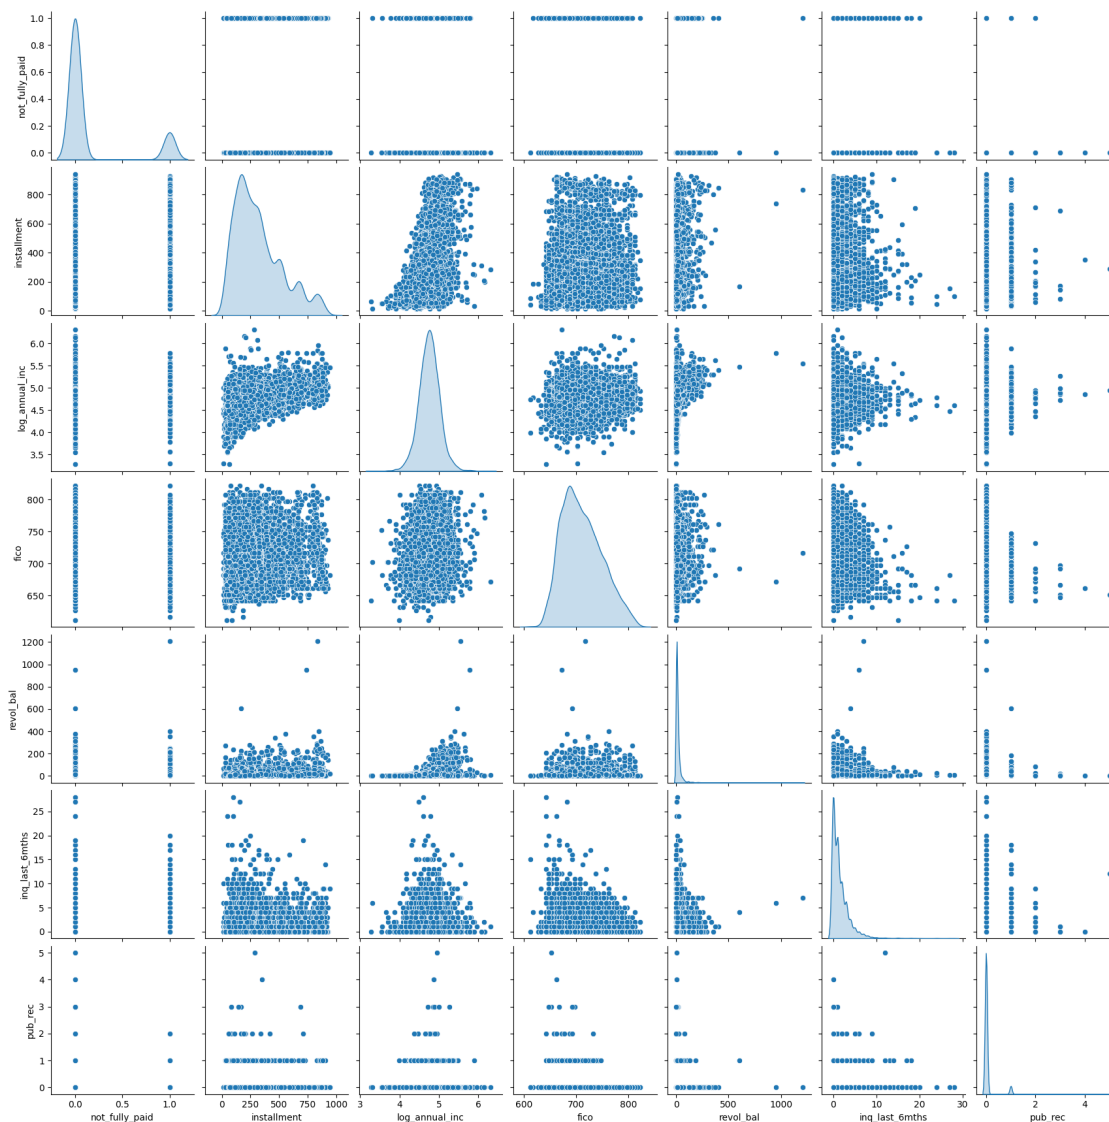
default_false = np.sum(loans_train['not_fully_paid'] == 0)
default_true = np.sum(loans_train['not_fully_paid'] == 1)

print(pd.Series({'0': default_false, '1': default_true}))
```

```
0    5585
1     1076
dtype: int64
```

```
[ ]: #use seaborn.pairplot to plot scatter plot for continuous variables
import seaborn as sns
sns.pairplot(loans_train,diag_kind='kde')
```

```
[ ]: <seaborn.axisgrid.PairGrid at 0x14e53a450>
```



```
[ ]: loans_train.corr()
```

```
[ ]:
      not_fully_paid  installment  log_annual_inc    fico \
not_fully_paid      1.000000      0.029430      -0.039698 -0.144830
installment         0.029430      1.000000      0.452113  0.094061
log_annual_inc     -0.039698      0.452113      1.000000  0.124615
fico               -0.144830      0.094061      0.124615  1.000000
revol_bal          0.040423      0.224112      0.362875 -0.010072
inq_last_6mths     0.141096     -0.005957      0.029720 -0.192888
pub_rec            0.045955     -0.028412      0.022748 -0.142629
```

	revol_bal	inq_last_6mths	pub_rec
not_fully_paid	0.040423	0.141096	0.045955
installment	0.224112	-0.005957	-0.028412
log_annual_inc	0.362875	0.029720	0.022748
fico	-0.010072	-0.192888	-0.142629
revol_bal	1.000000	0.026560	-0.027113
inq_last_6mths	0.026560	1.000000	0.066816
pub_rec	-0.027113	0.066816	1.000000

A baseline model can be a so-called “dummy” model, where the classifier predicts every new observation as the majority class. In our case, for a datapoint with any given features, the baseline model will always predict ‘no default’.

```
[ ]: # Accuracy of baseline model based on training data:
ACC = default_false/(default_false + default_true) #Accuracy = (TP+TN)/
      ↪ (TP+TN+FP+FN)
ACC
```

```
[ ]: 0.8384626932892959
```

Note: we want to create models that performs better than the baseline model.

```
[ ]: # EXERCISE: Compute accuracy of baseline on testing:
default_false_test = np.sum(loans_test['not_fully_paid'] == 0)
default_true_test = np.sum(loans_test['not_fully_paid'] == 1)
ACC_test = default_false_test/(default_false_test+default_true_test)
ACC_test
```

```
[ ]: 0.8441330998248686
```

```
[ ]: # EXERCISE: What are the TPR and FPR rates of the baseline model?

# True positive: the proportion of actual positives that are correctly
      ↪ identified as positive
# False positive: the proportion of actual negatives that are incorrectly
      ↪ identified as positive

TPR = 0/default_true_test # TPR = TP/P = TP/(TP+FN)
FPR = 0/default_false_test # FPR = FP/N = FP/(FP+TN)
print(TPR,FPR)
```

```
0.0 0.0
```

```
[ ]: # EXERCISE: What are the TNR and FNR rates of the baseline model?

# True negative: the proportion of actual negatives that are correctly
      ↪ identified as negative
# False negative: the proportion of actual positives that are incorrectly
      ↪ identified as negative
```

```
TNR = default_false_test/default_false_test #  $TNR = TN/N = TN/(TN+FP)$ 
FNR = default_true_test/default_true_test #  $FNR = FN/P = FN/(TP+FN)$ 
print(TNR,FNR)
```

1.0 1.0

1.1.5 1.5 Model Fitting (Logistic Regression)

Now we can use the statsmodels package to fit the training set to a logistic regression model

```
[ ]: import statsmodels.formula.api as smf
```

```
?smf.logit
```

Signature:

```
smf.logit(formula,
data, subset=None,
drop_cols=None,
*args,
**kwargs)
```

Docstring:

Create a Model from a formula and dataframe.

Parameters

formula : str or generic Formula object

The formula specifying the model.

data : array_like

The data for the model. See Notes.

subset : array_like

An array-like object of booleans, integers, or index values that indicate the subset of df to use in the model. Assumes df is a `pandas.DataFrame``.

drop_cols : array_like

Columns to drop from the design matrix. Cannot be used to drop terms involving categoricals.

*args

Additional positional argument that are passed to the model.

**kwargs

These are passed to the model with one exception. The `eval_env`` keyword is passed to patsy. It can be either a `:class:`patsy:patsy.EvalEnvironment`` object or an integer indicating the depth of the namespace to use. For example, the default `eval_env=0`` uses the calling namespace. If you wish to use a "clean" environment set `eval_env=-1``.

Returns

model

The model instance.

Notes

data must define `__getitem__` with the keys in the formula terms
args and kwargs are passed on to the model instantiation. E.g.,
a numpy structured or rec array, a dictionary, or a pandas DataFrame.

File:

/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-
packages/statsmodels/base/model.py

Type: method

```
[ ]: # Fit the logistic regression model

logreg = smf.logit(formula = 'not_fully_paid ~ installment + log_annual_inc +_
    fico + revol_bal + inq_last_6mths + pub_rec',
                    data = loans_train).fit(method = 'nag')

print(logreg.summary())
```

Optimization terminated successfully.

Current function value: 0.437108

Iterations: 4

Function evaluations: 5

Gradient evaluations: 5

Hessian evaluations: 4

Logit Regression Results

```
=====
Dep. Variable:          not_fully_paid    No. Observations:          6661
Model:                  Logit             Df Residuals:          6654
Method:                  MLE              Df Model:              6
Date:                   Wed, 13 Nov 2024   Pseudo R-squ.:           0.01154
Time:                   18:29:47          Log-Likelihood:          -2911.6
converged:               True             LL-Null:                 -2945.6
Covariance Type:         nonrobust         LLR p-value:             1.067e-12
=====
```

==

```
=====
              coef    std err          z      P>|z|      [0.025
0.975]
```

--

```
Intercept          -3.1e-06    0.901   -3.44e-06    1.000    -1.766
1.766
installment         0.0006     0.000    3.431     0.001     0.000
0.001
log_annual_inc     -1.513e-05   0.151    -0.000     1.000    -0.295
0.295
```

fico	-0.0026	0.001	-2.807	0.005	-0.004
-0.001					
revol_bal	0.0001	0.001	0.126	0.899	-0.002
0.002					
inq_last_6mths	2.443e-05	0.016	0.002	0.999	-0.031
0.031					
pub_rec	8.993e-07	0.125	7.19e-06	1.000	-0.245
0.245					

=====

==

1.1.6 1.6 Predictions

```
[ ]: # Example of prediction for a new observation

new_obs = pd.DataFrame(data = {'installment' : [366], 'log_annual_inc' : [4.
    ↪51], 'fico' : [682],
                                'revol_bal' : [7.53], 'inq_last_6mths' : [1],
    ↪'pub_rec' : [0]})

logreg.predict(new_obs)
```

```
[ ]: 0    0.172312
dtype: float64
```

```
[ ]: y_test = loans_test['not_fully_paid']

y_prob = logreg.predict(loans_test)
y_pred = pd.Series([1 if x > 0.5 else 0 for x in y_prob], index=y_prob.index)

# y_pred is the vector of probabilities as given by your model on the test set.
    ↪Values between 0 and 1.
# Remember,  $P(Y_i = 1) = 1/(1 + e^{-(b_0 + b_1*x_1 + b_2*x_2 + \dots)})$ 
```

1.1.7 1.7 Model Evaluation - Confusion Matrix

In order to evaluate the performance of our classification model, we can make use of confusion matrix to compute a variety of useful metrics

```
[ ]: from sklearn.metrics import confusion_matrix

?confusion_matrix
```

Signature:

```
confusion_matrix(
    y_true,
    y_pred,
    *,
    labels=None,
```



```
sample_weight=None,
```

```
normalize=None,
```

```
)
```

Docstring:

Compute confusion matrix to evaluate the accuracy of a classification.

By definition a confusion matrix C is such that $C_{i,j}$ is equal to the number of observations known to be in group i and predicted to be in group j .

Thus in binary classification, the count of true negatives is $C_{0,0}$, false negatives is $C_{1,0}$, true positives is $C_{1,1}$ and false positives is $C_{0,1}$.

Read more in the :ref:`User Guide <confusion_matrix>`.

Parameters

`y_true` : array-like of shape (n_samples,)
Ground truth (correct) target values.

`y_pred` : array-like of shape (n_samples,)
Estimated targets as returned by a classifier.

`labels` : array-like of shape (n_classes), default=None
List of labels to index the matrix. This may be used to reorder or select a subset of labels.
If ``None`` is given, those that appear at least once in ``y_true`` or ``y_pred`` are used in sorted order.

`sample_weight` : array-like of shape (n_samples,), default=None
Sample weights.

.. versionadded:: 0.18

`normalize` : {'true', 'pred', 'all'}, default=None
Normalizes confusion matrix over the true (rows), predicted (columns) conditions or all the population. If None, confusion matrix will not be normalized.

Returns

`C` : ndarray of shape (n_classes, n_classes)
Confusion matrix whose i -th row and j -th column entry indicates the number of samples with true label being i -th class

and predicted label being j-th class.

See Also

`ConfusionMatrixDisplay.from_estimator` : Plot the confusion matrix
given an estimator, the data, and the label.
`ConfusionMatrixDisplay.from_predictions` : Plot the confusion matrix
given the true and predicted labels.
`ConfusionMatrixDisplay` : Confusion Matrix visualization.

References

.. [1] `Wikipedia entry for the Confusion matrix
<https://en.wikipedia.org/wiki/Confusion_matrix>`_
(Wikipedia and other references may use a different
convention for axes).

Examples

```
>>> from sklearn.metrics import confusion_matrix
>>> y_true = [2, 0, 2, 2, 0, 1]
>>> y_pred = [0, 0, 2, 2, 0, 2]
>>> confusion_matrix(y_true, y_pred)
array([[2, 0, 0],
       [0, 0, 1],
       [1, 0, 2]])

>>> y_true = ["cat", "ant", "cat", "cat", "ant", "bird"]
>>> y_pred = ["ant", "ant", "cat", "cat", "ant", "cat"]
>>> confusion_matrix(y_true, y_pred, labels=["ant", "bird", "cat"])
array([[2, 0, 0],
       [0, 0, 1],
       [1, 0, 2]])
```

In the binary case, we can extract true positives, etc. as follows:

```
>>> tn, fp, fn, tp = confusion_matrix([0, 1, 0, 1], [1, 1, 1, 0]).ravel()
>>> (tn, fp, fn, tp)
(np.int64(0), np.int64(2), np.int64(1), np.int64(1))
```

File:

/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages/sklearn/metrics/_classification.py

Type: function

To remind you of what each element of the confusion matrix represents:

TN FP

FN TP

```
[ ]: cm = confusion_matrix(y_test, y_pred)
print ("Confusion Matrix : \n", cm)
```

```
Confusion Matrix :
[[2410    0]
 [ 445    0]]
```

```
[ ]: # compare the confusion matrix of the baseline model
baseline_model = [0]*loans_test.shape[0] # model of all zeros
confusion_matrix(loans_test['not_fully_paid'], baseline_model)
```

```
[ ]: array([[2410,    0],
          [ 445,    0]])
```

```
[ ]: print(cm.ravel())
```

```
[2410    0  445    0]
```

```
[ ]: # Accuracy
(cm.ravel()[0]+cm.ravel()[3])/sum(cm.ravel())
```

```
[ ]: 0.8441330998248686
```

Be careful about the definitions of FPR, TPR, recall, precision, sensitivity, specificity etc:
https://en.wikipedia.org/wiki/Sensitivity_and_specificity

```
[ ]: # EXERCISE: What is the True Positive Rate ?
TPR_logit = cm.ravel()[3]/(cm.ravel()[3]+cm.ravel()[2])
print('TPR is: %.4f' % TPR_logit)

# EXERCISE: What is the False Positive rate ?
FPR_logit = cm.ravel()[1]/(cm.ravel()[1]+cm.ravel()[0])
print('FPR is: %.4f' % FPR_logit)
```

```
TPR is: 0.0000
```

```
FPR is: 0.0000
```

```
[ ]: # Now, try threshold probability = 0.4
y_pred_40perc = pd.Series([1 if x > 0.4 else 0 for x in y_prob], index=y_prob.
    ↪index)
cm = confusion_matrix(y_test, y_pred_40perc)
print ("Confusion Matrix : \n", cm)
# EXERCISE: What is the Accuracy?
acc= (cm.ravel()[0]+cm.ravel()[3])/sum(cm.ravel())
print('Accuracy is: %.4f' %acc)

# EXERCISE: What is the True Positive Rate ?
TPR_logit = cm.ravel()[3]/(cm.ravel()[3]+cm.ravel()[2])
print('TPR is: %.4f' % TPR_logit)
```

```
# EXERCISE: What is the False Positive rate ?
FPR_logit = cm.ravel()[1]/(cm.ravel()[1]+cm.ravel()[0])
print('FPR is: %.4f' % FPR_logit)
```

Confusion Matrix :

```
[[2410    0]
 [ 445    0]]
```

Accuracy is: 0.8441

TPR is: 0.0000

FPR is: 0.0000

Take away: After running logistic regression, one can change the threshold probability. This will affect the predicted label. A smaller threshold will result in more observations being predicted as positive.

1.2 In-class activity: print out the confusion matrix when the threshold is 0.2

```
[ ]: y_pred_20perc = pd.Series([1 if x > 0.2 else 0 for x in y_prob], index=y_prob.
    ↪index)
cm = confusion_matrix(y_test, y_pred_20perc)
print("Confusion Matrix : \n", cm)

acc= (cm.ravel()[0]+cm.ravel()[3])/sum(cm.ravel())
print('Accuracy is: %.4f' %acc)

TPR_logit = cm.ravel()[3]/(cm.ravel()[3]+cm.ravel()[2])
print('TPR is: %.4f' % TPR_logit)

FPR_logit = cm.ravel()[1]/(cm.ravel()[1]+cm.ravel()[0])
print('FPR is: %.4f' % FPR_logit)
```

Confusion Matrix :

```
[[2313   97]
 [ 394   51]]
```

Accuracy is: 0.8280

TPR is: 0.1146

FPR is: 0.0402

1.3 2. Logistic Regression (Numerical + Categorical Variables)

1.3.1 1.1 Data Loading

```
[ ]: churn = pd.read_csv("customerchurn.csv")
churn.info()
churn.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7032 entries, 0 to 7031
```

Data columns (total 7 columns):

#	Column	Non-Null Count	Dtype
0	Churn	7032 non-null	int64
1	MonthlyCharges	7032 non-null	float64
2	SeniorCitizen	7032 non-null	int64
3	PaymentMethod	7032 non-null	object
4	InternetService	7032 non-null	object
5	tenure	7032 non-null	int64
6	Contract	7032 non-null	object

dtypes: float64(1), int64(3), object(3)

memory usage: 384.7+ KB

```
[ ]:   Churn  MonthlyCharges  SeniorCitizen  PaymentMethod  InternetService  \
0      0          29.85              0  Electronic check          DSL
1      0          56.95              0    Mailed check          DSL
2      1          53.85              0    Mailed check          DSL
3      0          42.30              0   Bank transfer          DSL
4      1          70.70              0  Electronic check  Fiber optic

      tenure      Contract
0      1  Month-to-month
1     34    One year
2      2  Month-to-month
3     45    One year
4      2  Month-to-month
```

1.3.2 1.2 Train Test Split

```
[ ]: churn_train, churn_test = train_test_split(churn, test_size=0.3,
      ↪random_state=88)
churn_train.shape, churn_test.shape
```

```
[ ]: ((4922, 7), (2110, 7))
```

1.3.3 1.3 Model Fitting and Variable Selection

```
[ ]: logreg = smf.logit(formula = 'Churn ~ MonthlyCharges + SeniorCitizen +
      ↪PaymentMethod + InternetService + tenure + Contract',
      data = churn_train).fit()
print(logreg.summary())
```

Optimization terminated successfully.

Current function value: 0.421913

Iterations 8

Logit Regression Results

```
=====
Dep. Variable:          Churn  No. Observations:          4922
```

```

Model:                Logit    Df Residuals:      4911
Method:                MLE     Df Model:           10
Date:                  Wed, 13 Nov 2024    Pseudo R-squ.:    0.2695
Time:                  18:29:48    Log-Likelihood:   -2076.7
converged:              True     LL-Null:         -2842.6
Covariance Type:        nonrobust    LLR p-value:      0.000

```

```

=====
=====

```

		coef	std err	z	P> z
[0.025	0.975]				

Intercept		-0.8982	0.213	-4.217	0.000
-1.316	-0.481				
PaymentMethod[T.Credit card]		-0.1066	0.134	-0.795	0.427
-0.369	0.156				
PaymentMethod[T.Electronic check]		0.4394	0.110	3.984	0.000
0.223	0.656				
PaymentMethod[T.Mailed check]		-0.0373	0.135	-0.277	0.782
-0.302	0.227				
InternetService[T.Fiber optic]		0.8217	0.154	5.336	0.000
0.520	1.123				
InternetService[T.No]		-0.7265	0.185	-3.925	0.000
-1.089	-0.364				
Contract[T.One year]		-0.8444	0.124	-6.815	0.000
-1.087	-0.602				
Contract[T.Two year]		-1.9124	0.226	-8.466	0.000
-2.355	-1.470				
MonthlyCharges		0.0067	0.004	1.857	0.063
-0.000	0.014				
SeniorCitizen		0.3852	0.097	3.981	0.000
0.196	0.575				
tenure		-0.0296	0.003	-11.559	0.000
-0.035	-0.025				

```

=====
=====

```

```

[ ]: # Let's remove MonthlyCharges
logreg2 = smf.logit(formula = 'Churn ~ SeniorCitizen + PaymentMethod +
↳InternetService + tenure + Contract',
                    data = churn_train).fit()
print(logreg2.summary())

```

```

Optimization terminated successfully.
Current function value: 0.422264
Iterations 8

```

Logit Regression Results

```

=====

```

Dep. Variable:	Churn	No. Observations:	4922
Model:	Logit	Df Residuals:	4912
Method:	MLE	Df Model:	9
Date:	Wed, 13 Nov 2024	Pseudo R-squ.:	0.2688
Time:	18:29:48	Log-Likelihood:	-2078.4
converged:	True	LL-Null:	-2842.6
Covariance Type:	nonrobust	LLR p-value:	0.000

```
=====
=====
```

	coef	std err	z	P> z
[0.025 0.975]				

Intercept	-0.5748	0.122	-4.724	0.000
-0.813 -0.336				
PaymentMethod[T.Credit card]	-0.1121	0.134	-0.836	0.403
-0.375 0.151				
PaymentMethod[T.Electronic check]	0.4389	0.110	3.982	0.000
0.223 0.655				
PaymentMethod[T.Mailed check]	-0.0463	0.135	-0.344	0.731
-0.311 0.218				
InternetService[T.Fiber optic]	1.0555	0.090	11.744	0.000
0.879 1.232				
InternetService[T.No]	-0.9328	0.147	-6.331	0.000
-1.222 -0.644				
Contract[T.One year]	-0.8110	0.122	-6.622	0.000
-1.051 -0.571				
Contract[T.Two year]	-1.8682	0.225	-8.321	0.000
-2.308 -1.428				
SeniorCitizen	0.3818	0.097	3.948	0.000
0.192 0.571				
tenure	-0.0280	0.002	-11.650	0.000
-0.033 -0.023				

```
=====
=====
```

```
[ ]: # Let's remove PaymentMethod Credit Card and PaymentMethodMailed check. How do
      ↳ we do this?
      # Create a new feature (dummy variable).

churn_train2 = churn_train.copy()
churn_train2['ElectronicCheck'] = (churn_train2['PaymentMethod'] == 'Electronic
      ↳ check').astype('int64')
churn_train2.drop(columns=['PaymentMethod'], inplace=True)

# Let's do the same for test set as well
```

```

churn_test2 = churn_test.copy()
churn_test2['ElectronicCheck'] = (churn_test2['PaymentMethod'] == 'Electronic_
↳check').astype('int64')
churn_test2.drop(columns=['PaymentMethod'], inplace=True)

churn_test2.head()

```

```

[ ]:      Churn  MonthlyCharges  SeniorCitizen  InternetService  tenure  \
2839      1          99.60           1      Fiber optic          4
4965      0          93.80           1      Fiber optic          13
6686      0          96.55           0      Fiber optic          5
4151      0          25.20           0              No          39
241       0         109.20           0      Fiber optic          72

```

```

      Contract  ElectronicCheck
2839  Month-to-month           1
4965  Month-to-month           0
6686  Month-to-month           1
4151      Two year           0
241    Two year           0

```

```

[ ]: logreg3 = smf.logit(formula = 'Churn ~ SeniorCitizen + ElectronicCheck +_
↳InternetService + tenure + Contract',
                        data = churn_train2).fit()
print(logreg3.summary())

```

Optimization terminated successfully.

Current function value: 0.422336

Iterations 8

Logit Regression Results

```

=====
Dep. Variable:      Churn      No. Observations:      4922
Model:              Logit      Df Residuals:          4914
Method:              MLE      Df Model:              7
Date:                Wed, 13 Nov 2024      Pseudo R-squ.:      0.2687
Time:                18:29:49      Log-Likelihood:      -2078.7
converged:           True      LL-Null:            -2842.6
Covariance Type:     nonrobust      LLR p-value:        0.000
=====

```

```

=====
                                coef      std err          z      P>|z|
-----
[0.025      0.975]
-----
Intercept                -0.6254      0.085      -7.363      0.000
-0.792      -0.459
InternetService[T.Fiber optic]      1.0549      0.089      11.817      0.000
0.880      1.230

```


InternetService[T.No]	-0.9313	0.146	-6.384	0.000
-1.217	-0.645			
Contract[T.One year]	-0.8125	0.122	-6.635	0.000
-1.053	-0.573			
Contract[T.Two year]	-1.8704	0.224	-8.332	0.000
-2.310	-1.430			
SeniorCitizen	0.3813	0.097	3.945	0.000
0.192	0.571			
ElectronicCheck	0.4901	0.080	6.098	0.000
0.333	0.648			
tenure	-0.0280	0.002	-11.928	0.000
-0.033	-0.023			

=====

=====

1.3.4 1.4 Predicting the Probability and Deciding the Threshold

```
[ ]: # 1. Predicting the probability of default
y_prob = logreg3.predict(churn_test2)

# 2. Determining the optimal threshold of the default probability
## price reduction = 800 and no price reduction = 1000
## 800 (1 - p/2) = 1000 (1 - p) -> p = 0.333 .
## The threshold of high churn-risk and low churn-risk is 0.333.

# 3. Predicting the label
y_pred = pd.Series([1 if x > 1/3 else 0 for x in y_prob], index=y_prob.index)
```

1.3.5 1.5 Confusion Matrix Given a Decision Threshold

```
[ ]: # Now we have our probabiltiy of default, we can construct a confusion matrix
      ↳based on decision tree threshold we have computed
# price reduction = 800 and no price reduction = 1000
# 800 (1 - p/2) = 1000 (1 - p) -> p = 0.333.
# High churn risk > 0.333.
# or high churn risk, the expected return of discount is higher than no
      ↳discount, b/c the prob of retention is lower for no discount

from sklearn.metrics import confusion_matrix
y_test = churn_test2['Churn']

cm = confusion_matrix(y_test, y_pred)
print ("Confusion Matrix : \n", cm)
```

Confusion Matrix :

```
[[1213  329]
 [ 154  414]]
```

change threshold

```
[ ]: threshold = 1/4
y_pred = pd.Series([1 if x > threshold else 0 for x in y_prob], index=y_prob.
    ↪index)
cm = confusion_matrix(y_test, y_pred)
print ("Confusion Matrix : \n", cm)
```

Confusion Matrix :

```
[[1082  460]
 [ 105  463]]
```

1.3.6 1.7 ROC Curves

The ROC curve plots the TPR and FPR for every break-even threshold p between 0.0 and 1.0

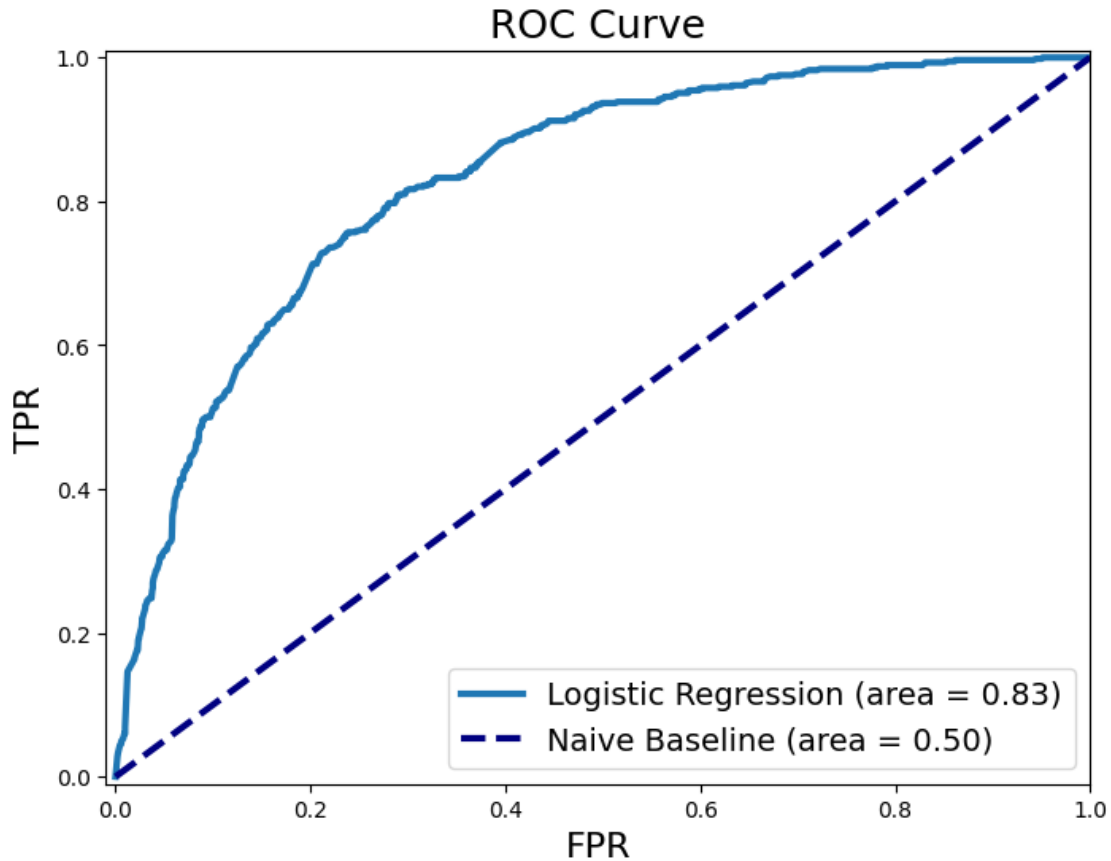
```
[ ]: y_train = churn_train2['Churn']
X_train = churn_train2.drop(['Churn'], axis=1)

y_test = churn_test2['Churn']
X_test = churn_test2.drop(['Churn'], axis=1)
```

```
[ ]: import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, auc

fpr, tpr, _ = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.title('ROC Curve', fontsize=18)
plt.xlabel('FPR', fontsize=16)
plt.ylabel('TPR', fontsize=16)
plt.xlim([-0.01, 1.00])
plt.ylim([-0.01, 1.01])
plt.plot(fpr, tpr, lw=3, label='Logistic Regression (area = {:.2f})'.
    ↪format(roc_auc))
plt.plot([0, 1], [0, 1], color='navy', lw=3, linestyle='--', label='Naive_
    ↪Baseline (area = 0.50)')
plt.legend(loc='lower right', fontsize=14)
plt.show()
```



A good model has a large area under the ROC curve, want it to get as close to the top left corner as possible.

1.4 2. Linear Discriminant Analysis

1.4.1 2.1 Preliminary: one-hot encoding

```
[ ]: from sklearn.preprocessing import OneHotEncoder
import pandas as pd

# Example data
df = pd.DataFrame({'Category': ['A', 'B', 'C', 'A']})

# OneHotEncoder = dummy variables
encoder = OneHotEncoder(sparse_output=False)
encoded_data = encoder.fit_transform(df[['Category']])
print(encoded_data)

# pd.get_dummies
dummies = pd.get_dummies(df, columns=['Category'])
```

```
print(dummies)

np.allclose(dummies, encoded_data)
```

```
[[1. 0. 0.]
 [0. 1. 0.]
 [0. 0. 1.]
 [1. 0. 0.]]
   Category_A  Category_B  Category_C
0         True         False         False
1         False          True         False
2         False         False          True
3         True          False         False
```

```
[ ]: True
```

```
[ ]: from sklearn.preprocessing import OneHotEncoder
      # initialize the OneHotEncoder
      drop_enc = OneHotEncoder(drop='first').
      ↪fit(X_train[['InternetService', 'Contract']])
      print(drop_enc.categories_)
```

```
[array(['DSL', 'Fiber optic', 'No'], dtype=object), array(['Month-to-month',
'One year', 'Two year'], dtype=object)]
```

```
[ ]: # Perform the transformation for both the training and the test set.

X_train_categorical = drop_enc.
    ↪transform(X_train[['InternetService', 'Contract']]).toarray()
X_train_numerical = X_train[['MonthlyCharges', 'SeniorCitizen', 'tenure']].values
# combine the numerical variables and the one-hot encoded categorical variables
X_train_transformed = np.concatenate((X_train_numerical, X_train_categorical),
    ↪axis = 1)

X_test_categorical = drop_enc.transform(X_test[['InternetService', 'Contract']]).
    ↪toarray()
X_test_numerical = X_test[['MonthlyCharges', 'SeniorCitizen', 'tenure']].values
X_test_transformed = np.concatenate((X_test_numerical, X_test_categorical), axis=
    ↪1)
```

1.4.2 2.2 Model Fitting and Prediction Making

```
[ ]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
      from sklearn.metrics import accuracy_score

      lda = LinearDiscriminantAnalysis()
      lda.fit(X_train_transformed, y_train)
```

```

y_prob_lda = lda.predict_proba(X_test_transformed)
y_pred_lda = pd.Series([1 if x > 1/3 else 0 for x in y_prob_lda[:,1]],
    ↪index=y_prob.index)

cm = confusion_matrix(y_test, y_pred_lda)
print ("Confusion Matrix: \n", cm)
print ("\nAccuracy:", accuracy_score(y_test, y_pred_lda))

```

Confusion Matrix:

```

[[1206  336]
 [ 163  405]]

```

Accuracy: 0.7635071090047393

1.4.3 2.3 Plot the ROC Curve

```

[ ]: fpr_lda, tpr_lda, _ = roc_curve(y_test, y_prob_lda[:,1])
roc_auc_lda = auc(fpr_lda, tpr_lda)

plt.figure(figsize=(8, 6))
plt.title('ROC Curve', fontsize=18)
plt.xlabel('FPR', fontsize=16)
plt.ylabel('TPR', fontsize=16)
plt.xlim([-0.01, 1.00])
plt.ylim([-0.01, 1.01])
plt.plot(fpr, tpr, lw=3, label='Logistic Regression (area = {:.2f})'.
    ↪format(roc_auc))
plt.plot(fpr_lda, tpr_lda, lw=3, label='LDA (area = {:.2f})'.
    ↪format(roc_auc_lda))
plt.plot([0, 1], [0, 1], color='navy', lw=3, linestyle='--', label='Naive
    ↪Baseline (area = 0.50)')
plt.legend(loc='lower right', fontsize=14)
plt.show()

```

